



# HYPER-CUBE

## HIGH-Dimensional Hypervisor Fuzzing




THE HIGHLY ANTICIPATED SEQUEL TO THE SMASH HIT



# CUBE 2: HYPERCUBE

There is more to fear than you can see.





# HYPER-CUBE

At a high-level HYPER-CUBE is a blackbox fuzzer designed to test hypervisors

## HYPER-CUBE components

- HYPER-CUBE OS - small, specialized OS
- TESSERACT - byte code interpreter for fuzzing
- A few accessory tools for working with TESSERACT



## Claimed contributions

1. Multi-dimensional, platform-independent fuzzing method
2. Fast hypervisor fuzzing
3. Custom OS for hypervisor fuzzing



## Prior Work

- VDF - AFL fuzzer for QEMU device emulators
- IOFUZZ - random writes to port mapped I/O
- Intel CHIPSEC - security tool that can fuzz hardware interfaces



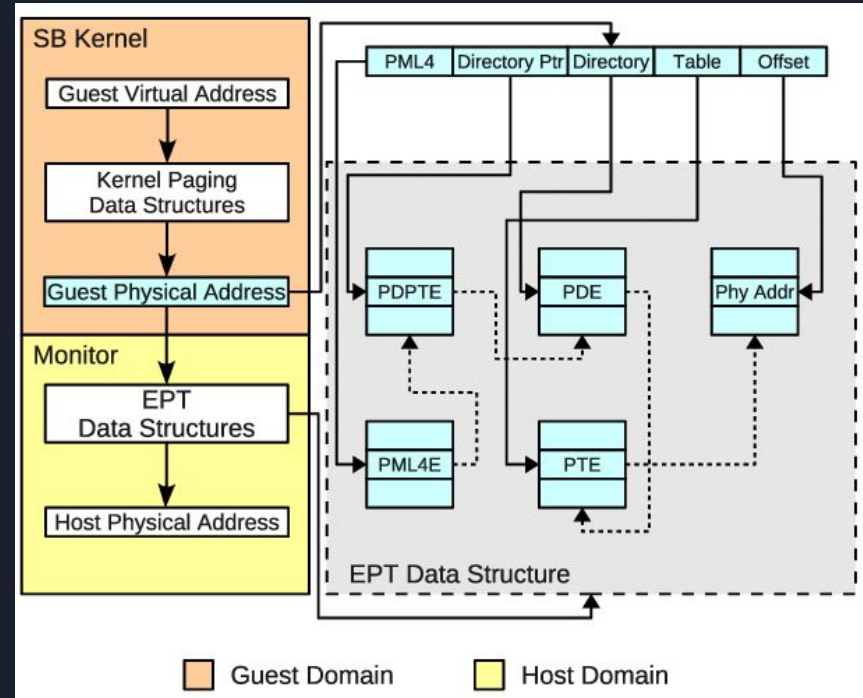
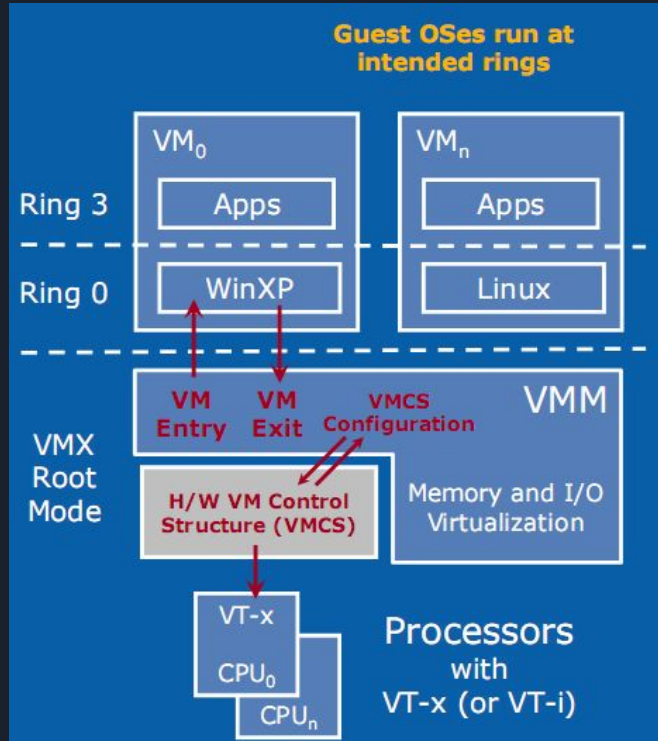
# X86 Virtualization Extensions

There are extensions to X86 that add hardware acceleration to virtualization, Intel VT-x and AMD-V.

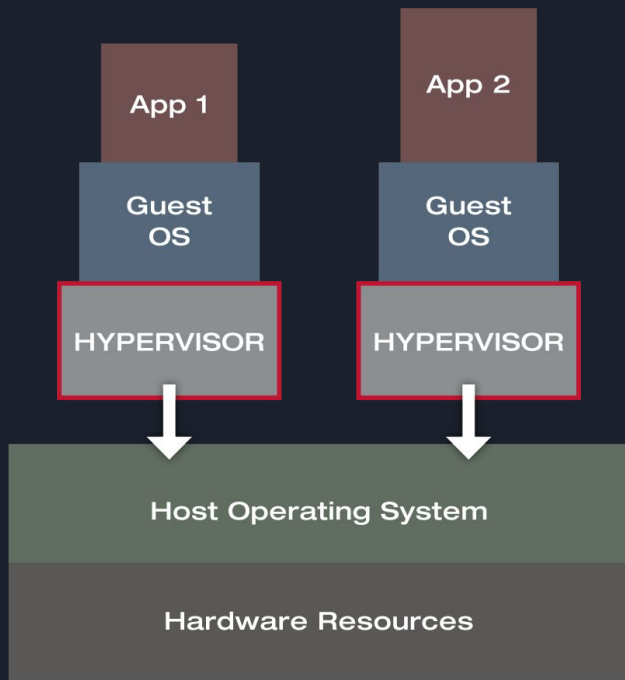
A few key extensions

- VMX - adds 13 new instructions, two new CPU modes
  - Root mode, used by the hypervisor
    - Can configure the Virtual Machine Control Structure (VMCS)
  - Non-root mode, used by guests
- Extended Page Tables (EPT)
  - Adds another level of page table translation
- Interrupt Virtualization
- I/O MMU Virtualization

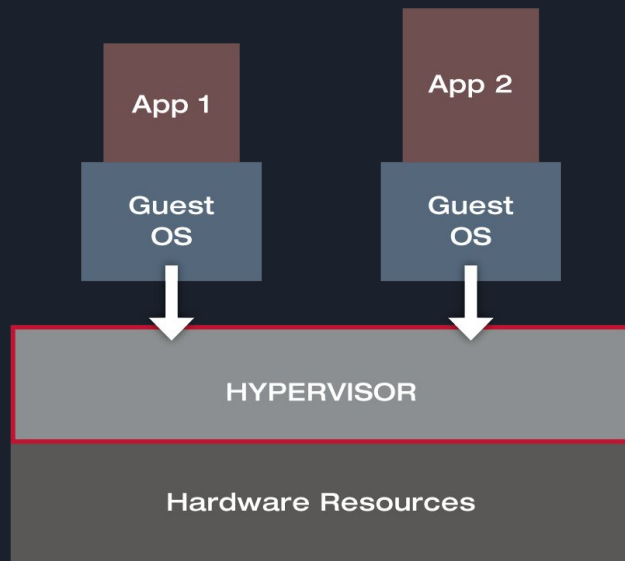
# X86 Virtualization Extensions



# Hypervisor Background



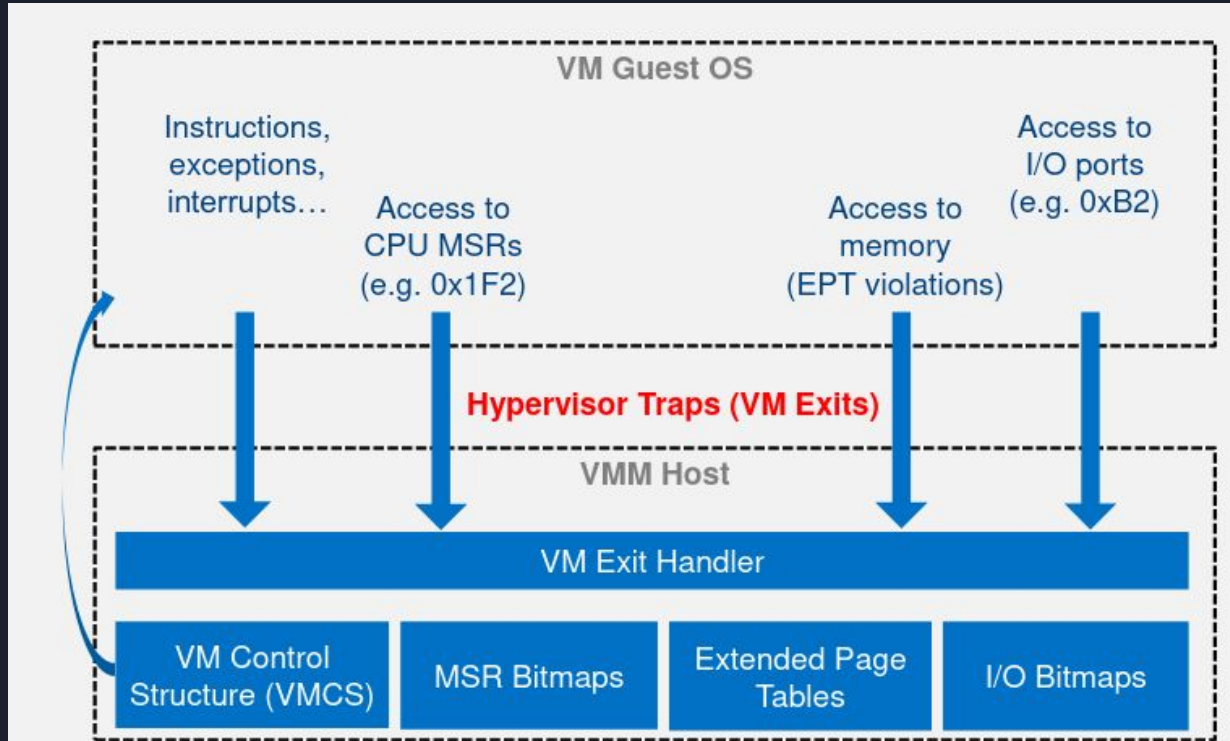
TYPE-2 HYPERVISOR



TYPE-1 HYPERVISOR



# VMExit





# VMExit

- Unconditional exit
  - VMX/SVM instructions
  - CPUID
  - GETSEC
  - INVD
  - XSETBV
- Conditional exit
  - CLTS
  - HLT
  - IN, INS/INSB/INSW/INSD, OUT, OUTS/OUTSB/OUTSW/OUTSD
  - INVLPG
  - INVPCID
  - LGDT, LIDT, LLDT, LTR, SGDT, SIDT, SLDT, STR
  - LMSW
  - MONITOR/MWAIT
  - MOV from CR3, CR8 / MOV to CR0, CR3, CR4, CR8
  - MOV DR
  - PAUSE
  - RDMSR/WRMSR
  - RDPMSR
  - RDRAND
  - RDTSCP
  - RSM
  - WBINVD
  - XRSTORS / XSAVES



# Design Goals

- X86 Hypervisor agnostic
- Blackbox fuzzing with high throughput
- High-dimensional
  - Interfaces
  - Operations



# OS/Fuzzing Architecture

## HYPER-CUBE OS

1. Multiboot 2 spec compliant, uses GRUB for booting OS
  - a. Can be boot via BIOS or UEFI
2. Simply memory management system
  - a. Uses a heap design and virtual addressing. Some virtual addresses map directly to physical address.
3. Interface enumerator
  - a. Finds most attached hardware interfaces
4. TESSERACT byte-code interpreter for fuzzing
5. Serial out interface
  - a. For outputting results



# Interfaces and enumeration

- Memory-mapped I/O (MMIO)
  - High precision event timers
  - Advanced programmable interrupt controller
  - PCI/PCIe devices
- Legacy Port I/O (PIO)
  - Programmable interrupt controller
  - ISA devices
  - Not exhaustive
- Direct Memory Access (DMA)
  - Not enumerated
- Hypercalls
  - Hardcoded, specific to each hypervisor



# TESSERACT

TESSERACT is a bytecode interpreter that generates calls to hypervisor interfaces.

- Is given a seed or generates ones. Seed is used to feed pseudorandom generator for creating a stream
- MMIO operations are given a `region_id` and offset. Is modulus size of respective range.
- Can be given dictionaries of data values and offsets



# TESSERACT

- `write_mmio(region id, offset, data)` writes a single word data to the address given by `region_id+offset`.
- `read_mmio(region id, offset)` reads a single word from the address given by `region_id+offset`.
- `xor_mmio(region id, offset, mask)` reads a single word from the given address, and writes it back after applying the given XOR mask.
- `bruteforce_mmio(region id, offset, data, num)` writes `num` consecutive data words to the given address.
- `memset_mmio(region id, offset, data, num)` writes the word data to `num` consecutive addresses, beginning at the given address.



# TESSERACT

- `writes_mmio(region id, offset, data, num)` same as `memset mmio`, however it uses a `rep` prefixed instruction to perform the task, testing instruction emulation.
- `reads_mmio(region id, offset, num)` same as `writes mmio`, but instead of writing data, it reads it.
- `mmio_write_scratch_ptr(region id, offset, scratch-id, scratch-offset)` writes a pointer to the given offset in the scratch area to the address in the given MMIO region.
- `*_io()` all opcodes accessing MMIO regions are implemented for I/O ports as well.
- `write_msr(msr num, mask)` writes to a MSR. This operation is limited to a list of  $\approx 240$  well-known MSRs. The mask is xored into the selected MSR.
- `hypercall(eax, ebx, ecx, edx, esi)` executes arbitrary hypercalls using the given registers as arguments. KVM
- `vmport(ecx, ebx)` executes arbitrary vmport hyper-calls with the registers set to the arguments to the hypervisor. VMware





# TESSERACT

- `writes_mmio(region id, offset, data, num)` same as `memset mmio`, however it uses a `rep` prefixed instruction to perform the task, testing instruction emulation.
- `reads_mmio(region id, offset, num)` same as `writes mmio`, but instead of writing data, it reads it.
- `mmio_write_scratch_ptr(region id, offset, scratch-id, scratch-offset)` writes a pointer to the given offset in the scratch area to the address in the given MMIO region.
- `*_io()` all opcodes accessing MMIO regions are implemented for I/O ports as well.
- `write_msr(msr num, mask)` writes to a MSR. This operation is limited to a list of  $\approx 240$  well-known MSRs. The mask is xored into the selected MSR.
- `hypercall(eax, ebx, ecx, edx, esi)` executes arbitrary hypercalls using the given registers as arguments. KVM
- `vmport(ecx, ebx)` executes arbitrary vmport hyper-calls with the registers set to the arguments to the hypervisor. VMware

# PRNG Stream

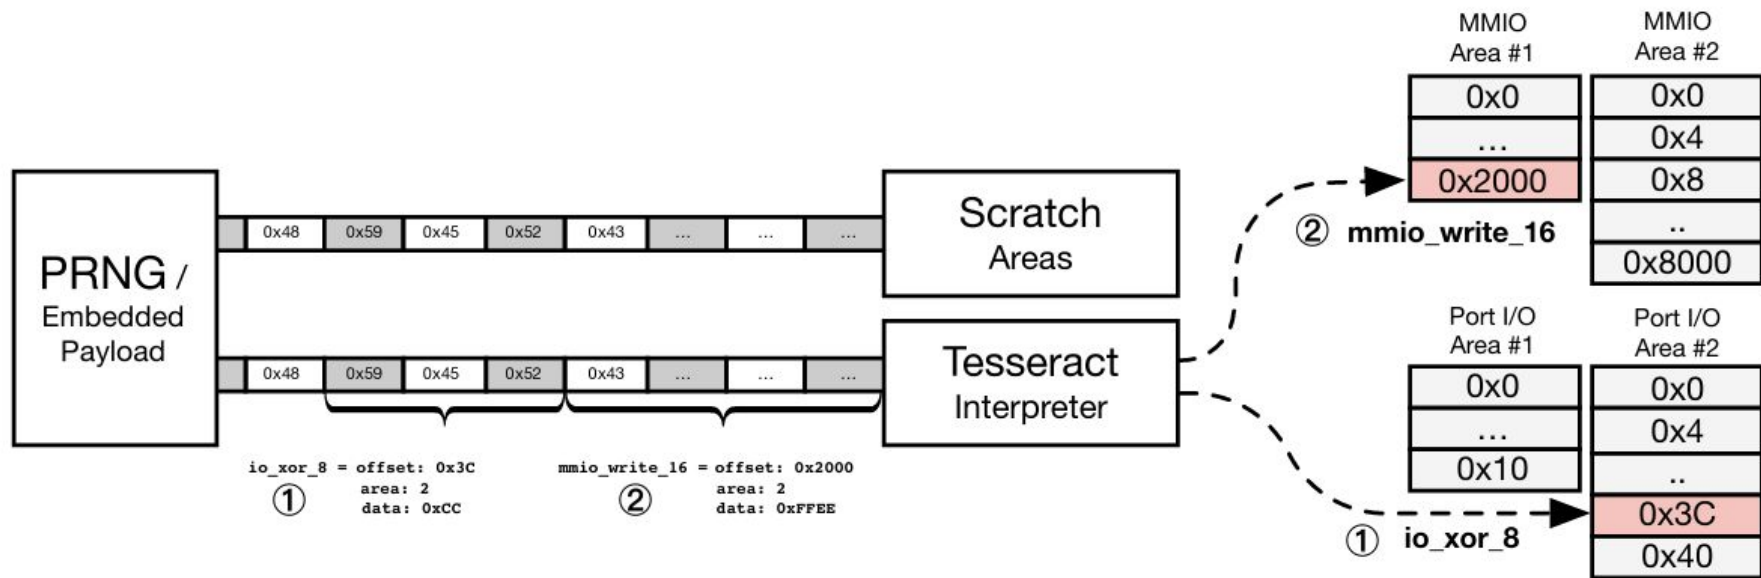
0120: 2fff 1c27 ab47 5700  
0128: adf2 3d60 092f 5488  
0130: ec2d 9d1a 029d 56fd  
0138: e0d1 a275 1f56 1d28  
0140: ea78 a2fa db07 d60d  
0148: 1288 3a5a 91f9 1756  
0150: 1cae 31ad 9b9c 938e  
0158: 2a33 f597 6615 e267  
0160: 0117 1f16 b440 8a86  
0168: 9154 5b55 e4ca 9e3d  
0170: 9d19 ae79 efac e500  
0178: 8cdf 8c00 9a83 df76  
0180: 91fe d779 026c 2e2b  
0188: 9137 1ef8 eea3 d29c  
0190: 1789 5938 a36f 718a  
0198: 81e4 678c 20f5 fa0b  
01a0: 774d 07f1 cee3 62bc  
01a8: d845 bc86 7631 6eac

**Robust  
Interpretation**



# Opcode Handler

**vmport**(0xbd4,0x10ea)  
**memset\_io**(0x426,0xce0,0x9dc,0xca8)  
  
**writes\_mmio**(0xec8,0xad,0x10ac,0x7e9)  
**bruteforce\_mmio**(0xce4,0xdfa,0xe31,0x322)  
  
**writes\_io**(0x4bb,0xb8,0xeb1,0x401)  
  
**memset\_mmio**(0x128,0xa73,0x2b3,0xa84)  
**read\_mmio**(0xbf3,0x907)  
  
**bruteforce\_io**(0x5c4,0x49a,0x94f,0xb1c)  
  
**xor\_mmio**(0x54b,0xa00,0xb51)

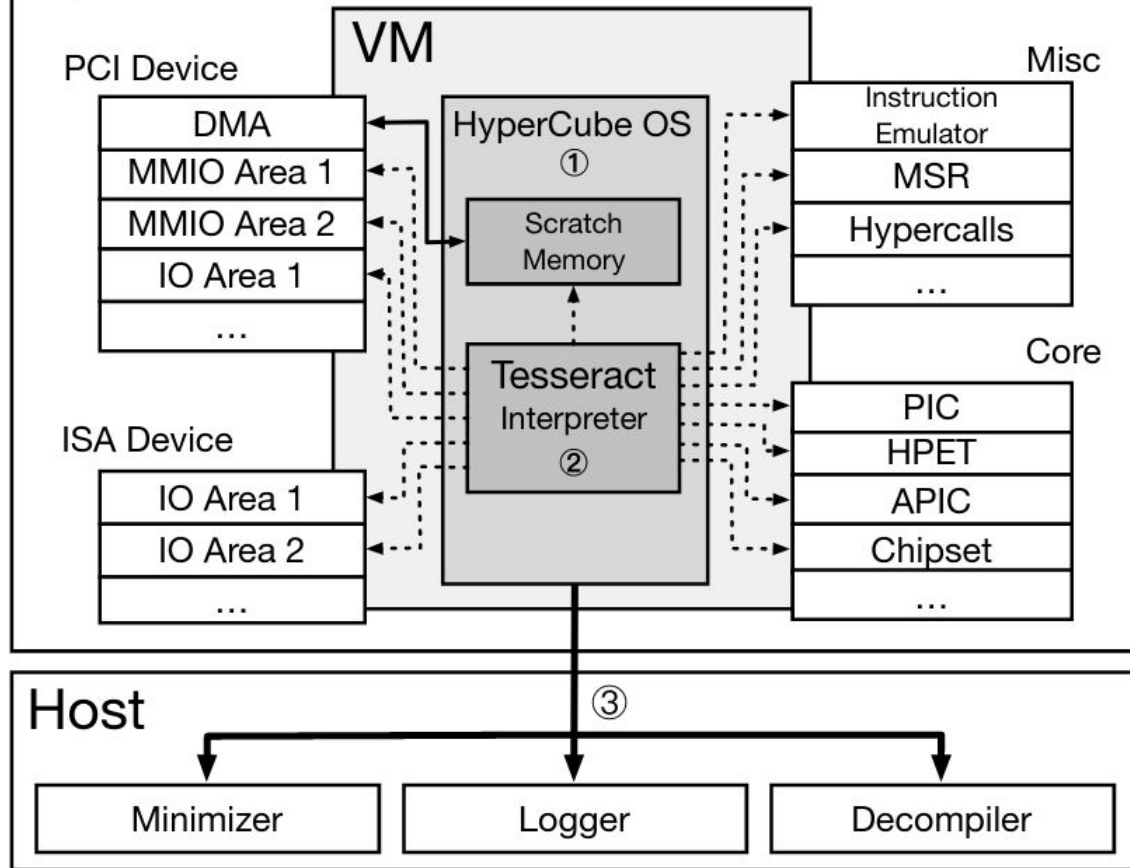




# External Tools

- Logger
  - Accepts serial communication from virtual machine
- Results minimizer
  - Tries to reduce triggering input stream down to smallest possible trigger
- Decompiler
  - For interrupting results

# Hypervisor





# Tested Hypervisors

- FreeBSD bhyve (12.0-RELEASE)
- VirtualBox (5.1.37\_Ubuntu r122592)
- Parallels Desktop (14.1.3)
- KVM/QEMU (4.0.1-rc4)
- Intel ACRN (29360 Build)
- VMware Fusion (11.0.3)



# Results

Found 55 bugs, 43 CVEs

- Assert Failures 25
- Null-Pointer Dereferences 13
- Memory-Corruptions 8
- Div-By-Zero (FP Exceptions) 5
- Deadlocks 4



# HYPER-CUBE vs VDF

Fuzzed 15 Device Emulators

## HYPER-CUBE

- 13/15 More Coverage
- 9/15 Crashes
- 10 Minutes Each

## VDF

- 2/15 More coverage
- 4/15 Crashes
- ~ 60 days each





# Limitations

1. Restrictions on some Type-1 hypervisors
2. Paravirtualized components need hand written interfaces
3. Black-box fuzzing