



Format

- Welcome
- Discuss Confidentiality
- Issue Clearing (if required)
- Meaningful Update & Follow-Up
- Review Next Papers
- Presentation
- Feedback



REVelry Branch History Injection

Cross-Privilege Spectre-v2 Attacks &
Mitigations

About Me | Jay Warne

Currently

- DARPA research
 - Side-channels
 - Processors
 - Hypervisors
 - Rowhammer Style Attacks
 - Program Analysis Methods
- Project Work
 - TL for Enclosed Training Environment
 - PO for Videographic Data Analysis
- Advisory
 - Product Development
 - Product Direction
 - Expert Reviewer

Previously

- Ran a Security Ops Team
- Occasional RE & Red Team
- Field Forensic Analysis & Tool Deployment

Things I Like

- Alpine Ski Racing
- Ski/ Alpine Mountaineering
- Ice/ Rock Climbing
- Surfing
- Backpacking






Branch History Injection (BHI) – tl;dr

- Branch Target Injection (BTI) aka Spectre-v2 exploits branch mispredictions
- Hardware and software mitigations exist, but to limited benefit
- BHI works with latest hardware mitigations in place, leaking kernel from userland

PoC Code Available At:

[http://vusec.net/projects/
bhi-spectre-bhb](http://vusec.net/projects/bhi-spectre-bhb)





What We Will Cover

- Related Concepts
 - Branch Prediction
 - Branch Target Injection (BTI)
 - uArch Leaking & Disclosure Gadgets
 - Transient Execution
- Mitigations for BTI
 - Retpoline (Intel/ AMD)
 - Workaround_1 (ARM)
 - IBPB
 - STIBP
 - eIBRS
 - FEAT_CSV2
 - CET-IBT
- Branch History Injection
 - Defeating Countermeasures
 - Attack Surface
 - Requirements
- Exploitation
 - Primitives
 - Exploitation w/ eBPF
 - Exploitation w/o eBPF
- Results

Related Concepts

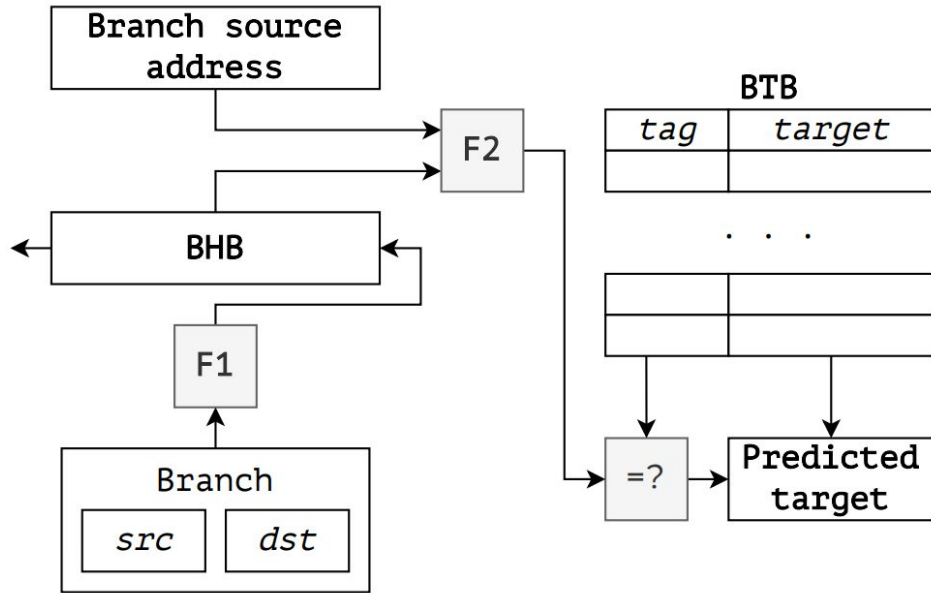




Related Concepts

- Branch Prediction
- Branch Target Injection (BTI)
- uArch Leaking & Disclosure Gadgets
- Transient Execution
- Branch Prediction Unit (BPU)
- Predicts the targets of branches based on previous behavior to speculatively execute on upcoming instructions
- Direct Branches
- Indirect Branches

Related Concepts



- Branch Prediction Unit (BPU)
- Predicts the targets of branches based on previous behavior to speculatively execute on upcoming instructions
- Direct Branches
- Indirect Branches



Related Concepts

- Branch Prediction
 - Branch Target Injection (BTI)
 - uArch Leaking & Disclosure Gadgets
 - Transient Execution
- Was one of the Spectre gadgets released
 - Processors try to predict targets
 - Mistrained processor speculatively fetches the wrong thing
 - If the processor selects the wrong branch, we can arbitrarily execute code on the victim



Related Concepts

- Branch Prediction
 - Branch Target Injection (BTI)
 - uArch Leaking & Disclosure Gadgets
 - Transient Execution
- This work is predicated in part on knowledge of side channel leaks and disclose
 - “Fallout” toy example



Snippets

```
volatile char* lut = (volatile char *)mmap(NULL, STRIDE*256, PROT_READ|PROT_WRITE, MAP_SHARED|MAP_ANONYMOUS,
volatile char* v_pg = (volatile char *)mmap(NULL, PAGE_SIZE*50, PROT_READ|PROT_WRITE, MAP_SHARED|MAP_ANONYMOUS,
volatile char* a_pg = NULL; //(volatile char *)mmap(NULL, PAGE_SIZE, PROT_NONE, MAP_PRIVATE|MAP_ANONYMOUS, -
if (v_pg == MAP_FAILED || lut == MAP_FAILED || a_pg == MAP_FAILED) {
    perror("MMAP ERROR BOIS");
    return 1;
}
mprotect((char *)a_pg, PAGE_SIZE, PROT_NONE);
memset((void *)v_pg, 1, PAGE_SIZE*50);
memset((void *)lut, 4, STRIDE*256);

warmup_memory();
warmup_memory();
warmup_memory();
warmup_memory();
```

Snippets

```
clflush(&v_pg[offset]);
```

```
for (int i = 0; i < table_entries; i++) {  
    clflush(lut + STRIDE * i);  
    mfence();  
}
```

```
cpuid(0);  
unsigned status = _xbegin();  
if(_XBEGIN_STARTED == status) {  
    for (int x = 0; x < 45; x++) {  
        v_pg[offset + x * STRIDE] = 42;  
    }  
    *(lut + STRIDE * a_pg[offset]);  
    /*(lut + STRIDE * 5);  
    _xend();  
}
```

```
for (int i = 0; i < table_entries; ++i) {  
    struct results & res = rarray[i];  
    uint32_t aux_val = 0;  
    mfence();  
    cpuid(0);  
    uint64_t perf = read_perf(0);  
    uint64_t time = rdtscp(aux_val);  
    lut[i * STRIDE];  
    mfence();  
    //cpuid(0);  
    uint64_t time2 = rdtscp(aux_val);  
    uint64_t perf2 = read_perf(0);  
    res.dt = time2 - time;  
    res.dp = perf2 - perf;  
}
```

Mitigations for BTI





Mitigations for BTI

- Retpoline (Intel/ AMD)
- Workaround_1 (ARM)
- IBPB
- STIBP
- eIBRS
- FEAT_CSV2
- CET-IBT



Mitigations for BTI

- Retpoline (Intel/ AMD)
- Workaround_1 (ARM)
- IBPB
- STIBP
- eIBRS
- FEAT_CSV2
- CET-IBT
- Software Mitigation
- Code sequence that converts indirect branches to `ret` instructions (intel) or `lfence; jmp` sequences (AMD)
- Intel ensures use of the Return Stack Buffer (RSB) instead of BTB
- AMD ensures that the load retires before the jump, reducing the transient execution window
- Anything that would mispeculate just loops until the address is resolved
- No indirect branch prediction is performed



Mitigations for BTI

- Retpoline (Intel/ AMD)
- Workaround_1 (ARM)
- IBPB
- STIBP
- eIBRS
- FEAT_CSV2
- CET-IBT
- Software Mitigation
- Attempts to invalidate branch predictor entries by
 - executing `BPIALL` on mode switch
 - Disabling and re-enabling the MMU



Mitigations for BTI

- Retpoline (Intel/ AMD)
- Workaround_1 (ARM)
- IBPB
- STIBP
- eIBRS
- FEAT_CSV2
- CET-IBT
- Hardware Mitigation
- Indirect Branch Prediction Barrier
- Previous branches don't affect the execution of subsequent branches



Mitigations for BTI

- Retpoline (Intel/ AMD)
- Workaround_1 (ARM)
- IBPB
- STIBP
- eIBRS
- FEAT_CSV2
- CET-IBT
- Hardware Mitigation
- Single Thread Indirect Branch Predictors
- Restrict sharing of branch prediction state across hyper-thread and cores



Mitigations for BTI

- Retpoline (Intel/ AMD)
- Workaround_1 (ARM)
- IBPB
- STIBP
- eIBRS
- FEAT_CSV2
- CET-IBT
- Hardware Mitigation
- Indirect Branch Restricted Speculation
- Allegedly prevents indirect branch prediction from being controlled by lower-privileged parts



Mitigations for BTI

- Retpoline (Intel/ AMD)
- Workaround_1 (ARM)
- IBPB
- STIBP
- eIBRS
- FEAT_CSV2
- CET-IBT
- Hardware Mitigation
- ARM's version of: Indirect Branch Restricted Speculation
- Supports two different solutions:
 - Prevent speculative control of indirect branch targets from different contexts
 - Enforce the above guarantees for the software context



Mitigations for BTI

- Retpoline (Intel/ AMD)
 - Workaround_1 (ARM)
 - IBPB
 - STIBP
 - eIBRS
 - FEAT_CSV2
 - CET-IBT
- Hardware Mitigation
 - Control-Flow Enforcement Technology – Indirect Branch Tracking
 - Prevents speculative and architectural execution of all indirect branch targets that don't start with `ENDBR32/64`

Branch History Injection





Defeating Countermeasures

- Author's thought was that indirect branch prediction is complex, so corners may have been cut to not lose performance
- Syscall

Defeating Countermeasures

- Author's thought was that indirect branch prediction is complex, so corners may have been cut to not lose performance
- Syscall

Algorithm 1 Experiment to verify BHB isolation between user and kernel mode. For the training phase, $n = 1$ is usually sufficient on the tested CPUs in Table 2.

```
1: for  $i \leftarrow 1, n$  do                                     ▷ Training
2:   set_history( $H_{load}$ )
3:   syscall  $\rightarrow$  load(mem)           ▷  $H_{load} \rightarrow load(mem)$ 
4:   set_history( $H_{dummy}$ )
5:   syscall  $\rightarrow$  dummy                 ▷  $H_{dummy} \rightarrow dummy$ 
6: end for
7: flush(mem)                                     ▷ Remove mem from cache
8: set_history( $H_{load}$ )
9: syscall  $\rightarrow$  dummy           ▷ Misprediction to load(mem)?
10: reload(mem)
```



Defeating Countermeasures

- Author's thought was that indirect branch prediction is complex, so corners may have been cut to not lose performance
- Syscall
- Success proves that the BHB is not isolated
- BHI can overcome defenses to divert execution, but **only to valid kernel targets**
- Their testing showed success on hypervisors as well

Attack Surface

- Out-of-place BTI is possible since BTB collisions can be created only by controlling BHB
- Out-of-place BTI not possible on ARM
- In-place BTI appears to be possible everywhere

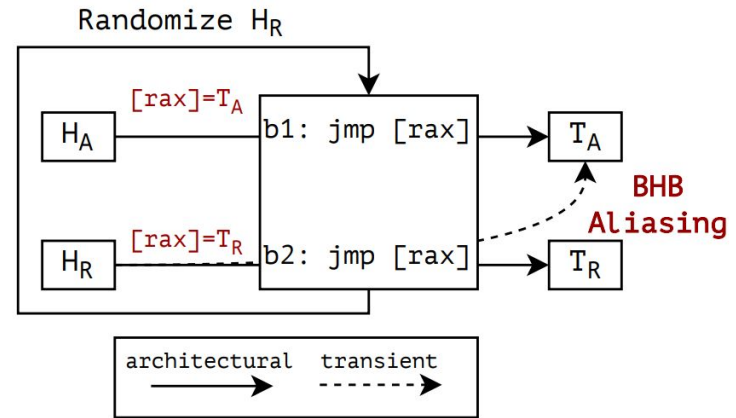


Figure 4: Brute-force approach to find colliding BHB values.

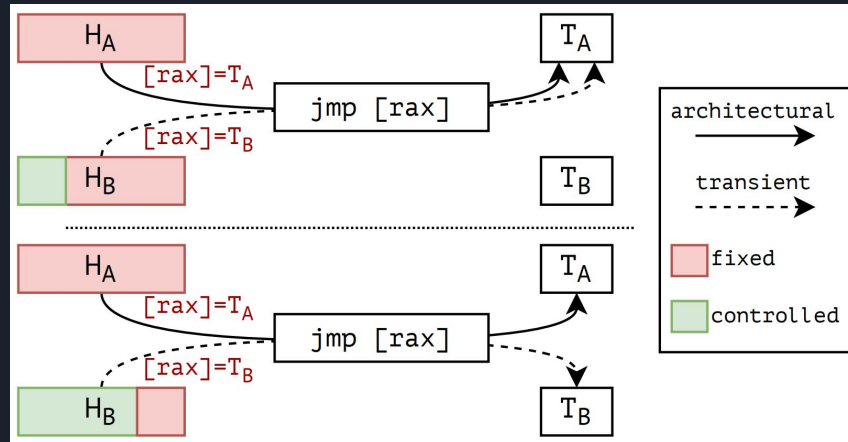


Requirements

- BHB Size
- History Brute-forcing
- History Controllability

Requirements

- BHB Size
- History Brute-forcing
- History Controllability



- Knowing the size informs how many indirect branches needed to cause mis-speculation
- Test for size loop
 - Given: Two colliding histories: H_A , H_B
 - Step 1: Alter oldest branch of H_B
 - Step 2: Test if H_B causes misprediction
 - Exit: When enough old branches of H_B have been altered to no longer cause misprediction

Requirements

- BHB Size
- History Brute-forcing
- History Controllability

- Must be able to generate matching BTB tags to cause mispeculation
- Brute Force Method
 - Given: Fixed History and Target: H_A , T_A
 - Step 1: Generate Random History (H_R) ending at some Random Target (T_R)
 - Step 2: Test if H_R causes misprediction to T_A
 - Exit: When H_R causes BTB tag to match T_A and cause mispeculation

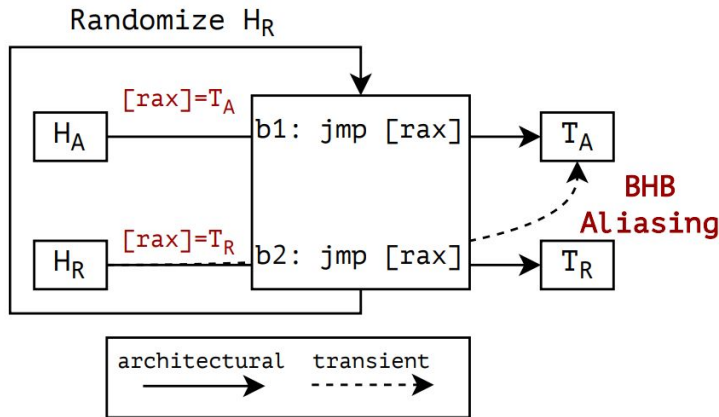
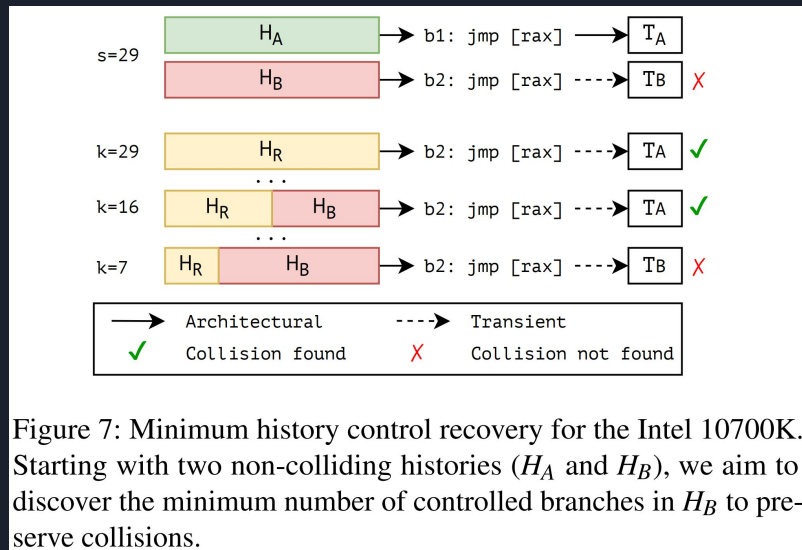


Figure 4: Brute-force approach to find colliding BHB values.

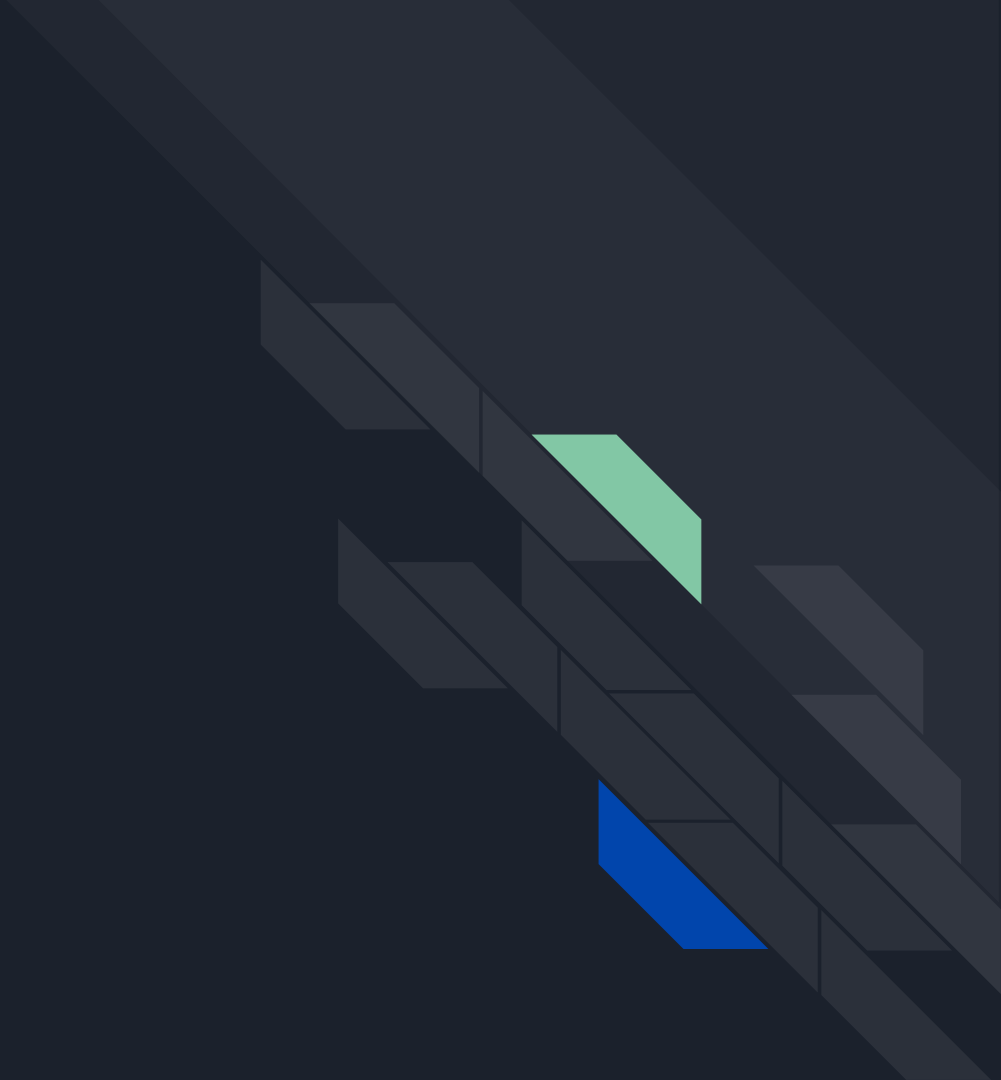
Requirements

- BHB Size
- History Brute-forcing
- History Controllability

- How many branches must be controlled to generate arbitrary BTB tags
- Using H_R from History Brute-forcing
- Fix branches closest to Target jump leading to misprediction to a known bad history, H_B until misprediction fails



Exploitation



Full Attack

1. Disclosure Gadget (F+R) from eBPF
2. TAG Generated by Disclosure Gadget is kernel-tagged
3. Execute Jump Chain from 'Requirements' to cause BHB collision before syscall
4. Tag collision with disclosure gadget instead of syscall handler
5. "Transient Type Confusion" Gadget leaks into r12, recovered by F+R covert channel

Listing 2 JIT'ted code for the eBPF program serving as disclosure gadget.

```
1  push    rbp
2  mov     rbp, rsp
3  ;load er_buf base address
4  movabs  rsi, 0xffffc900028ff110
5  ;rdi+0x18 = &pt_regs.r12 transiently
6  ;        = &bpf_sock architecturally
7  mov     rax, QWORD PTR [rdi+0x18]
8  test    rax, rax
9  je      fail
10 ;Dereference of user r12 value transiently
11 mov     eax, DWORD PTR [rax+0x14]
12 ;extract the byte to leak
13 and     rax, 0xff
14 shl     rax, 0xc
15 add     rsi, rax
16 ;maccess(er_buf[byte_to_leak*0x1000])
17 mov     rsi, QWORD PTR [rsi+0x0]
18 fail:
19 xor     eax, eax
20 leave
21 ret
```